

⏪

🔍 Find by title

Microsoft Copilot Studio
documentation

▾ Overview

What is Microsoft Copilot Studio

Choose a harness

Access standard harness agents and
agent flows

What's new

> Get started

> Billing and licensing

> Agents

> Workflows

> Security, data residency, and
compliance

> Administration

Find help and support

> Troubleshooting

> Languages and locales

> Respond to personal data requests

> Responsible AI

↓ Download PDF

Learn > Microsoft Copilot Studio >


 **Ask Learn** ⋮

Copilot Studio overview

✦ Summarize this article for me

Walk through this page, talk about what Copilot Studio is in general, and mention that we have a limited implementation of it as part of the pilot.

Microsoft Copilot Studio is a graphical, low-code studio for building and managing AI-powered agents and workflows. Build agents and workflows, connect them to your organization's data and systems, and publish them to the channels where your users already work.



Untitled Agent

COPILOT STUDIO

Building Agents at First Hawaiian Bank

Instructions



John Scalla
Officer, Systems Analyst

Talk about how Sean and John's team put together a great overview of using Co-pilot Studio at First Hawaiian Bank, and everything they need to know starts on this page here.

Then walk through the page and show everything that's there, commenting on different things.

This guide will help you get started quickly on building a Copilot Studio agent for use within First Hawaiian Bank, so that your agent can be reviewed and promoted.

You build in a dedicated, sealed environment called **FHB Agent Dev**. Inside this environment, you have full freedom to create, test, and break things. You cannot publish to the organization, share org-wide, export, or promote. When an agent is ready, you **request promotion**, and a controlled review-and-release process takes it from there.

THE THREE RULES

1. Create a solution first. Never build an agent outside a solution.
2. One agent per solution. If an agent needs help from another agent, you either build a child agent inside it or connect a separate agent. Don't put two standalone agents in one solution. If you think you have a reason to, speak to a Copilot Architect/Engineer for advice.
3. Set your solution as your preferred one, and reset it whenever you switch to a different agent.

Why the rules?

If you create an agent without first creating and selecting a solution, Copilot Studio silently drops it into the Default solution.

An agent in the Default solution can never be promoted to Test or Production. It is stuck in Dev forever.

There is no "fix it later" that isn't rework. This is why Rule #1 exists.

Learn the Flow

Follow these steps in order to ensure your agent is built correctly and can be promoted for review and release. Each step links to a detailed guide:

Agent Build Steps

Talk about why these steps are important. And for starting out, they'll only have to go through the first three steps. I haven't gone through all six, I just went through the point to develop my first application.



1. Before You Build

Confirm you're in the right place; the rules and the naming...



2. Create Your Solution First

The required first step, every time.



3. Create Your Agent Inside the Solution

Plus the 10-second check that proves it worked



4. Switching Between Solutions

The habit that keeps your work from landing in the wrong place.



5. One Agent per Solution + Connecting Agents

How agents call other agents, and the rules that keep that clean.



6. When You're Ready to Ship

You request promotion; you don't do it yourself.

Before You Build

 John Scalla
Officer, Systems Analyst

This is how to confirm you’re in the right place. Also show the right place from Copilot Studio for my particular area.

Navigation

Create Your Solution First

One Agent per Solution, and How Agents Work Together

Building Agents at First Hawaiian Bank

Create Your Agent Inside the Solution

When You’re Ready to Ship

Switching Between Solutions

Make Sure You're in “FHB Agent Dev”

Everything you do must happen in FHB-Agent-Dev. If you build in any other environment, your work won't be governed and can't be promoted.

- Open [Copilot Studio](#)
- Look at the environment name at the bottom of the left navigation. It must read FHB-Agent-Dev.
- If it doesn't, select the environment picker and switch to FHB-Agent-Dev.

Do this **before** you create your agent. New agent → new solution.


Step 1

Open the Solution Explorer

1. In the left navigation, select the ellipsis (...) near the bottom.
2. Under Explore, select Solutions.

Tip: The Solution Explorer is your starting point for every agent build. Always begin here to ensure your work is governed and promotable.

Talk about the importance of developing their apps within a solution. The reason for that is if they don't develop it within a solution, it will never be able to be promoted and shared.




Step 2

Start a New Solution

Select + New solution at the top of the Solution Explorer.

Creating your solution first is critical. If you skip this, your agent will be placed in the Default solution, which cannot be promoted. Always create your solution before your agent.



Step 3

New solution

Display name *

Current preferred solution
FHB Agent - AI Assessor
[Manage](#)

Create plan from a solution
Select an existing solution, and AI agents will generate a detailed plan showing how that solution supports your business goal

Here's a look at the different solutions that have been put together so far. Walk through the list.

Unmanaged
Managed
All

Display name	Name	Created	Version	Publisher	Solution check	Source control status
FHB Agent - Security Test	FHBAgentSecurityTest	3 days ago	1.0.0.0	First Hawaiian Bank	Hasn't been run	Not Connected to Git
FHB Agent - AI Assessor	AIApplicationAssessor	5 days ago	1.0.0.0	First Hawaiian Bank	Results as of 8:14 AM, 8/5/26	Not Connected to Git
FHB Agent - AIWG AI Use Case Assessment	FHBAgentAIWGAIUs...	6 days ago	1.0.0.0	First Hawaiian Bank	Hasn't been run	Not Connected to Git
FHB Agent - EISTR Controls Testing	FHBAgentEISTRCont...	6 days ago	1.0.0.0	First Hawaiian Bank	Hasn't been run	Not Connected to Git
FHB Agent - FDIC Docs Reviewer	FHBAgentFDICDocsR...	2 weeks ago	1.0.0.0	First Hawaiian Bank	Hasn't been run	Not Connected to Git
FHB Agent - Team News	FHBAgentTeamNews	2 weeks ago	1.0.0.0	First Hawaiian Bank	Hasn't been run	Not Connected to Git
FHB Agent - Daily Brief	FHBAgentDailyBrief	2 weeks ago	1.0.0.0	First Hawaiian Bank	Hasn't been run	Not Connected to Git
FHB Agent - Shawn Test 2	FHBAgentShawnTest2	3 weeks ago	1.0.0.0	First Hawaiian Bank	Hasn't been run	Not Connected to Git
FHB Agent - FDIC	FHBAgentFDIC	3 weeks ago	1.0.0.0	First Hawaiian Bank	Hasn't been run	Not Connected to Git
FHB Agent - Test Policy Agent	FHBAgentTestPolicy...	3 weeks ago	1.0.0.0	First Hawaiian Bank	Hasn't been run	Not Connected to Git
FHB Agent - RMG Budget Helper	FHBAgentRMGBudg...	3 weeks ago	1.0.0.0	First Hawaiian Bank	Hasn't been run	Not Connected to Git

Building Agents at First Hawaiian Bank

When You're Ready to Ship

Switching Between Solutions

Before You Build

One Agent per Solution, and How
Agents Work Together

Create Your Solution First

You've [created and preferred your solution](#). Now the agent will land in it automatically.

There are two safe ways to get an agent into your Solution.

Path A (Recommended)

If you're first starting out, always go with the recommended solution. It'll be a lot easier. In this case, we'll follow path A. I haven't explored the other paths.

Create Your Agent Inside the Preferred Solution

- With your solution already set as preferred, simply create the agent, and it will be placed correctly.
- Step 1: Go to the [Copilot Studio Home](#) page.
- Step 2: Describe what you want your agent to do in plain language, or choose to configure it manually.
- Step 3: Copilot Studio creates the agent inside your preferred solution.

Since your solution is preferred, you can just create the agent normally, and it will go to the right place.

REMINDER — Setting a preferred solution means that any assets you create or save will automatically be saved in that solution. Always switch your preferred solution when working within another solution.

Path B

Voila! Now we're ready to build agents. Cover some of the other things that are on the left-hand side of the menu.

Hey James, ready to automate?

Optimize your business processes Uses Copilot Credits

Create Planner tasks when messages are flagged in Teams|



Agent GitHub Copilot

Create an agent to take actions and answer questions.



Workflow GitHub Copilot

Automate multi-step processes with triggers, actions, and AI steps.

Other ways to build
Create agents that extend Microsoft 365 Copilot, or build standard agents and automated flows with rule-based orchestration. →



Agents

1 item

Search by name

Name		Status	Owner	Last m
	AI Security Assessor	... Draft	James Jaworski	Aug 7

Here is an agent that I built, and we'll take a look at that.

Tell the story behind it.

Show each of the supporting documents.

Back in April, or maybe even it was the end of March, I was trying to think of what were some of the risk factors for agentic AI applications. That was part of developing the AI security roadmap. I also had a full AI security program outlined as well.

Risk Factors for Agentic AI Applications

1. Autonomy and Human Oversight

The more independently an agent can act — making decisions, executing multi-step plans, and taking real-world actions without human checkpoints — the higher the risk. A read-only research agent is very different from one that can send emails, execute code, or move money. Key considerations include the degree of human-in-the-loop control, the availability of approval gates at critical decision points, and whether the agent can be paused or overridden in real time.

2. Scope of Access and Permissions

What systems, data, and tools can the agent reach? An agent with access to production databases, financial systems, or PII carries far more risk than one sandboxed to public information. This includes both the breadth of access (how many systems) and depth (what it can do within each system). Least-privilege principles are essential here.

3. Reversibility of Actions

Can the agent's actions be undone? Sending a Slack message is hard to unsend; deleting a database table is much harder to reverse than writing a draft document. Irreversible actions in high-stakes domains (finance, healthcare, legal) amplify risk considerably. Agents that only draft or recommend — rather than execute — carry inherently lower risk.

4. Task Complexity and Ambiguity

When tasks are well-defined and narrow, agents are more predictable. As tasks become open-ended — "research this topic and take appropriate action" — the space for unintended behavior grows. Ambiguous instructions leave room for the agent to interpret goals in ways that diverge from user intent. Clear task boundaries and well-specified success criteria reduce this risk.

Risk Factor Rubric

The following section provides detailed scoring guidance for each of the 12 risk factors. For each factor, five distinct levels are described from Very Low (1) to Very High (5). Use the level descriptions and example characteristics to determine the appropriate score for the application under review.

From there, I try to come up with a rubric for each of these factors so it could determine what level of maturity or risk within each of the areas.

1. Autonomy and Human Oversight

Weight: 1.5

The degree to which the agent can act independently—making decisions, executing multi-step plans, and taking real-world actions without human checkpoints. This is the single most defining characteristic of agentic AI risk.

Score	Level Description	Characteristics
1 Very Low	Fully human-controlled; agent only provides suggestions or drafts	<i>Agent outputs are advisory only</i>
2 Low	Every action requires explicit human approval before execution	<i>Human approval gate on every action</i>
3 Medium	Human approves critical actions; routine actions auto-execute	<i>Routine actions auto-execute; critical ones need approval</i>
4 High	Agent acts autonomously with periodic human review after the fact	<i>Autonomous with periodic after-the-fact review</i>
5 Very High	Fully autonomous operation with no human checkpoints	<i>Fully autonomous with no human checkpoints</i>

2. Scope of Access and Permissions

Weight: 1.4

What systems, data, and tools the agent can reach. This includes both the breadth of access

A	B	C	D	E
Risk Factor Weight Rationale				
#	Risk Factor	Weight	Rationale	
1	Autonomy and Human Oversight	1.5	Highest weight. The degree of autonomy is the defining characteristic of agentic AI risk — an agent that can act without human approval has fundamentally different risk than one that only suggests.	
2	Scope of Access and Permissions	1.4	High weight. Access scope directly determines what damage is possible. An agent can only be as dangerous as the systems it can reach.	
3	Reversibility of Actions	1.3	High weight. Irreversible actions in high-stakes domains can cause permanent harm. The ability to undo mistakes is a critical safety property.	
4	Task Complexity and Ambiguity	1.0	Standard weight. While ambiguity increases unpredictability, clear task boundaries are a solvable design problem with less inherent danger than access or autonomy.	
5	Tool and Chain-of-Action Risk	1.1	Step also offers an intervention opportunity.	
6	Data Sensitivity and Privacy	1.4	High weight. Data sensitivity carries regulatory, legal, and reputational risk that extends well beyond operational concerns.	
7	Transparency and Explainability	0.8	Lower weight. Important for governance and accountability, but transparency is a control mechanism rather than a direct risk driver.	
8	Reliability and Failure Modes	1.1	Slightly elevated. Failure handling determines whether problems stay small or cascade, making it a critical amplifier of other risks.	

Then I built a spreadsheet to kind of come up with a weighting scheme for these factors. And it's all kind of arbitrary, although I did run it by clod to get clods and put on what I was putting together.



AI Security Assessor



Instructions



AI Security Architect — Agentic AI Risk Assessment Prompt

You are an internal AI Security Architect at a regulated financial institution.

PURPOSE

Conduct a structured risk assessment of proposed agentic AI applications using the this 12-factor Agentic AI Risk Assessment model. Your role is advisory only: identify risk factors, explain why they matter, and recommend mitigations. You do NOT approve or reject use cases.

BOUNDARIES

- Decision-support only; no legal or compliance sign-off
- Assume current-state controls only (not future intentions)
- Do not ask for or accept PII, NPPI, credentials, or customer data
- If sensitive examples are needed, instruct users to use placeholders

INTERACTION STYLE

- Ask one question at a time
- Briefly explain why each question matters (one sentence)
- Use minimal follow-ups only if needed for clarity
- Ensure all 12 risk factors are evaluated before producing the report

INTAKE (START EVERY SESSION)

Ask for:

1. Owner
2. Business Unit
3. 1–2 sentence description of what the AI application is intended to do

RISK FACTORS AND WEIGHTS

Score each factor from 1 (Very Low) to 5 (Very High), using conservative judgment.

So I took all that and then wrapped it up in terms of a co-pilot studio app. At the time I had a 30-day free license to Copilot Studio. And so I put together an app just to test out what Copilot Studio was and then kind of left it. But then, when we started moving forward with Copilot Studio within the bank, I was actually able to go back and extend my 30-day trial and was able to copy out what I had for this one particular app, and that's kind of where we are now. So we just started with it. I think it was like Thursday of last week because I had difficulty using the FHB platform, but I'll give you a demo how it works.

After the demo, we could go through and talk about what the instructions are for the agent.

Model

Skills

Define behaviors through

Tools

Connect the agent to o

Knowledge

Provide trusted contex

Search all web

Connected agents

Collaborate across age

Memory Preview

Remember interaction
improved results.